

CSCI 5451

Week 9 Notes

Professor Ellen Gethner

Shortest Path Problems, continued from last week

How to Grow a Tree

- **Definition.** Let G be a simple graph and T a subgraph of G that is a tree.

Shortest Path Problems, continued from last week

How to Grow a Tree


- ▶ **Definition.** Let G be a simple graph and T a subgraph of G that is a tree.
- ▶ A **frontier edge** of T is an edge $uv \in E(G)$ such that $u \in V(T)$ and $v \notin V(T)$.

Shortest Path Problems, continued from last week


How to Grow a Tree

- ▶ **Definition.** Let G be a simple graph and T a subgraph of G that is a tree.
- ▶ A **frontier edge** of T is an edge $uv \in E(G)$ such that $u \in V(T)$ and $v \notin V(T)$.
- ▶ **Example.** $G = K_4$

Example: frontier edges




Example: frontier edges




Tree T , so far, is the single green vertex.

Example: frontier edges




The current frontier edges are **red**.

Example: frontier edges




Choose a frontier edge and update the **tree**.

Example: frontier edges




Now update the **frontier edges**.

Example: frontier edges




Choose any frontier edge and update the **tree**.

Example: frontier edges



And update the **frontier edges** again.

Example: frontier edges



And finally, choose one last frontier edge and update the **tree**.
Why are we done?

Tree Growing Algorithm

- ▶ **Input:** A connected graph G and starting vertex $v \in V(G)$
- ▶ **Output:** A spanning tree T of G and a vertex labeling of $V(G)$.
- ▶ **initialize** T as vertex v
- ▶ **initialize** set of frontier edges for T as **empty**
- ▶ **write** label 0 on vertex v
- ▶ **initialize** label counter $i = 1$
- ▶ **while** T does not yet span G
- ▶ **update** the set of frontier edges for T
- ▶ **let** e be the frontier edge for T of highest priority
- ▶ **let** w be the unlabeled endpoint of e
- ▶ **add** edge e (and vertex w) to tree T
- ▶ **write** label i on vertex w
- ▶ $i = i + 1$
- ▶ **return** T with its vertex labeling

An application of tree growing

- ▶ Back at the ranch we were interested in finding the shortest path from some root to each other vertex in a graph.

An application of tree growing

- ▶ Back at the ranch we were interested in finding the shortest path from some root to each other vertex in a graph.
- ▶ Two things that are missing so far from our tree growing algorithm are

An application of tree growing

- ▶ Back at the ranch we were interested in finding the shortest path from some root to each other vertex in a graph.
- ▶ Two things that are missing so far from our tree growing algorithm are
 1. the fact that the edges are weighted, and

An application of tree growing

- ▶ Back at the ranch we were interested in finding the shortest path from some root to each other vertex in a graph.
- ▶ Two things that are missing so far from our tree growing algorithm are
 1. the fact that the edges are weighted, and
 2. a way to assign priorities to the frontier edges.

An application of tree growing

- ▶ Back at the ranch we were interested in finding the shortest path from some root to each other vertex in a graph.
- ▶ Two things that are missing so far from our tree growing algorithm are
 1. the fact that the edges are weighted, and
 2. a way to assign priorities to the frontier edges.
- ▶ So next let's use the weights of the edges to prioritize frontier edges in a way that makes sense for an eventual shortest path algorithm.

Prioritizing Frontier Edges

- ▶ Let G be a simple undirected weighted graph and suppose T is a subgraph of G that is a tree.

Prioritizing Frontier Edges

- ▶ Let G be a simple undirected weighted graph and suppose T is a subgraph of G that is a tree.
- ▶ For each frontier edge e of T , define the **priority of** e by $P(e) = d(s, x) + \omega(e)$, where x is the labeled endpoint of e ;

Prioritizing Frontier Edges

- ▶ Let G be a simple undirected weighted graph and suppose T is a subgraph of G that is a tree.
- ▶ For each frontier edge e of T , define the **priority of** e by $P(e) = d(s, x) + \omega(e)$, where x is the labeled endpoint of e ;
- ▶ recall that $d(s, x)$ is the distance from vertex s to vertex x , and $\omega(e)$ is the weight of edge e (see week 8's notes).

Prioritizing Frontier Edges

- ▶ Let G be a simple undirected weighted graph and suppose T is a subgraph of G that is a tree.
- ▶ For each frontier edge e of T , define the **priority of** e by $P(e) = d(s, x) + \omega(e)$, where x is the labeled endpoint of e ;
- ▶ recall that $d(s, x)$ is the distance from vertex s to vertex x , and $\omega(e)$ is the weight of edge e (see week 8's notes).
- ▶ Now that we know how to prioritize the frontier edges, we have all of the tools needed to describe **Dijkstra's Single Source Shortest Path Algorithm**.

Algorithm Dijkstra(G, s)

- **Input.** A weighted, connected, undirected graph G whose edge weights are non-negative; a source vertex $s \in V(G)$.

Algorithm Dijkstra(G, s)

- ▶ **Input.** A weighted, connected, undirected graph G whose edge weights are non-negative; a source vertex $s \in V(G)$.
- ▶ **Output.** A spanning tree T (the “Dijkstra Tree”) of G , rooted at vertex s such that

Algorithm Dijkstra(G, s)

- ▶ **Input.** A weighted, connected, undirected graph G whose edge weights are non-negative; a source vertex $s \in V(G)$.
- ▶ **Output.** A spanning tree T (the “Dijkstra Tree”) of G , rooted at vertex s such that
 1. the unique path from s to each vertex $v \in V(G)$ in T is a shortest path from s to v in G , and

Algorithm Dijkstra(G, s)

- ▶ **Input.** A weighted, connected, undirected graph G whose edge weights are non-negative; a source vertex $s \in V(G)$.
- ▶ **Output.** A spanning tree T (the “Dijkstra Tree”) of G , rooted at vertex s such that
 1. the unique path from s to each vertex $v \in V(G)$ in T is a shortest path from s to v in G , and
 2. a vertex labeling giving the distance from s to each vertex.

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
- ▶ **update** the frontier edges for T

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest
 - ▶ **let** x be the unlabeled endpoint of e and y be the labeled endpoint of e

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest
 - ▶ **let** x be the unlabeled endpoint of e and y be the labeled endpoint of e
 - ▶ **add** edge e (and vertex y) to tree T

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest
 - ▶ **let** x be the unlabeled endpoint of e and y be the labeled endpoint of e
 - ▶ **add** edge e (and vertex y) to tree T
 - ▶ $dist[y] = P(e)$

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest
 - ▶ **let** x be the unlabeled endpoint of e and y be the labeled endpoint of e
 - ▶ **add** edge e (and vertex y) to tree T
 - ▶ $dist[y] = P(e)$
 - ▶ **write** label $dist(y)$ on vertex y .

Algorithm Dijkstra(G, s), continued

- ▶ **initialize** the Dijkstra tree T as the given source vertex s
- ▶ **initialize** the set of frontier edges as **empty**
- ▶ $dist[s] = 0$
- ▶ **write** label 0 on vertex s
- ▶ **while** the Dijkstra tree T does not yet span G
 - ▶ **update** the frontier edges for T
 - ▶ **for** each frontier edge $e = xy$ for T
 - ▶ **let** x be the labeled endpoint of e
 - ▶ **let** y be the unlabeled endpoint of e
 - ▶ **set** $P(e) = dist[x] + \omega(e)$.
 - ▶ **let** $e = xy$ be a frontier edge of T for which $P(e)$ is smallest
 - ▶ **let** x be the unlabeled endpoint of e and y be the labeled endpoint of e
 - ▶ **add** edge e (and vertex y) to tree T
 - ▶ $dist[y] = P(e)$
 - ▶ **write** label $dist(y)$ on vertex y .
- ▶ **return** Dijkstra tree T and its vertex labels.

Thoughts about the algorithm

- ▶ **Remark.** Each time a new vertex y (together with its associated edge) is added to the Dijkstra tree, no new computations are needed to update the frontier edges;

Thoughts about the algorithm

- ▶ **Remark.** Each time a new vertex y (together with its associated edge) is added to the Dijkstra tree, no new computations are needed to update the frontier edges;
- ▶ simply delete frontier edges that include the endpoint y .

Thoughts about the algorithm

- ▶ **Remark.** Each time a new vertex y (together with its associated edge) is added to the Dijkstra tree, no new computations are needed to update the frontier edges;
- ▶ simply delete frontier edges that include the endpoint y .
- ▶ Then recompute the priorities of the remaining edges.

Thoughts about the algorithm

- ▶ **Remark.** Each time a new vertex y (together with its associated edge) is added to the Dijkstra tree, no new computations are needed to update the frontier edges;
- ▶ simply delete frontier edges that include the endpoint y .
- ▶ Then recompute the priorities of the remaining edges.
- ▶ **Question.** Is Dijkstra's algorithm correct?

Proof of Correctness of Dijkstra's Algorithm

- **Theorem.** Let T_j be the Dijkstra tree after j iterations of Algorithm Dijkstra on connect graph G for $j = 0, 1, 2, \dots, |V(G)| - 1$. Then for each $v \in V(T_j)$, the unique path from s to v in T_j is a shortest path from s to v in G .

Proof of Correctness of Dijkstra's Algorithm

- ▶ **Theorem.** Let T_j be the Dijkstra tree after j iterations of Algorithm Dijkstra on connect graph G for $j = 0, 1, 2, \dots, |V(G)| - 1$. Then for each $v \in V(T_j)$, the unique path from s to v in T_j is a shortest path from s to v in G .
- ▶ **Proof by induction.**

Proof of Correctness of Dijkstra's Algorithm

- ▶ **Theorem.** Let T_j be the Dijkstra tree after j iterations of Algorithm Dijkstra on connect graph G for $j = 0, 1, 2, \dots, |V(G)| - 1$. Then for each $v \in V(T_j)$, the unique path from s to v in T_j is a shortest path from s to v in G .
- ▶ **Proof by induction.**
- ▶ **Base Cases.**

Proof of Correctness of Dijkstra's Algorithm

- ▶ **Theorem.** Let T_j be the Dijkstra tree after j iterations of Algorithm Dijkstra on connect graph G for $j = 0, 1, 2, \dots, |V(G)| - 1$. Then for each $v \in V(T_j)$, the unique path from s to v in T_j is a shortest path from s to v in G .
- ▶ **Proof by induction.**
- ▶ **Base Cases.**
 - ▶ The claim is true for T_0 . Why?

Proof of Correctness of Dijkstra's Algorithm

- ▶ **Theorem.** Let T_j be the Dijkstra tree after j iterations of Algorithm Dijkstra on connect graph G for $j = 0, 1, 2, \dots, |V(G)| - 1$. Then for each $v \in V(T_j)$, the unique path from s to v in T_j is a shortest path from s to v in G .
- ▶ **Proof by induction.**
- ▶ **Base Cases.**
 - ▶ The claim is true for T_0 . Why?
 - ▶ What must you check to know that the claim is true for T_1 ?

Proof of correctness of Dijkstra, continued

- ▶ **Induction Hypothesis.** Assume for some fixed j such that $0 \leq j \leq |V(G)| - 1$ that T_j satisfies the claim (Theorem on previous slide).

Proof of correctness of Dijkstra, continued

- ▶ **Induction Hypothesis.** Assume for some fixed j such that $0 \leq j \leq |V(G)| - 1$ that T_j satisfies the claim (Theorem on previous slide).
- ▶ **Inductive Step.** Let $e \in E(G)$ with $e = xy$ and suppose $x \in V(T_j)$ and $y \notin V(T_j)$.

Proof of correctness of Dijkstra, continued

- ▶ **Induction Hypothesis.** Assume for some fixed j such that $0 \leq j \leq |V(G)| - 1$ that T_j satisfies the claim (Theorem on previous slide).
- ▶ **Inductive Step.** Let $e \in E(G)$ with $e = xy$ and suppose $x \in V(T_j)$ and $y \notin V(T_j)$.
- ▶ In particular, e is a frontier edge, x is labeled, and y is unlabeled.

Proof of correctness of Dijkstra, continued

- ▶ **Induction Hypothesis.** Assume for some fixed j such that $0 \leq j \leq |V(G)| - 1$ that T_j satisfies the claim (Theorem on previous slide).
- ▶ **Inductive Step.** Let $e \in E(G)$ with $e = xy$ and suppose $x \in V(T_j)$ and $y \notin V(T_j)$.
- ▶ In particular, e is a frontier edge, x is labeled, and y is unlabeled.
- ▶ Suppose e is the frontier edge added to T_j in the $(j + 1)$ st iteration of Algorithm Dijkstra.

Proof of correctness of Dijkstra, continued

- ▶ **Induction Hypothesis.** Assume for some fixed j such that $0 \leq j \leq |V(G)| - 1$ that T_j satisfies the claim (Theorem on previous slide).
- ▶ **Inductive Step.** Let $e \in E(G)$ with $e = xy$ and suppose $x \in V(T_j)$ and $y \notin V(T_j)$.
- ▶ In particular, e is a frontier edge, x is labeled, and y is unlabeled.
- ▶ Suppose e is the frontier edge added to T_j in the $(j + 1)$ st iteration of Algorithm Dijkstra.
- ▶ Since y is the only new vertex in T_{j+1} , it suffices to show that the path from s to y in T_{j+1} is a shortest path from s to y in G .

Proof of correctness of Dijkstra, continued

- ▶ Let Q be the (unique) path from s to y in T_{j+1} .

Proof of correctness of Dijkstra, continued

- ▶ Let Q be the (unique) path from s to y in T_{j+1} .
- ▶ **Observation.** The path Q contains a path from s to x in T_j and adds vertex y (and hence edge e) to form a path from s to y .

Proof of correctness of Dijkstra, continued

- ▶ Let Q be the (unique) path from s to y in T_{j+1} .
- ▶ **Observation.** The path Q contains a path from s to x in T_j and adds vertex y (and hence edge e) to form a path from s to y .
- ▶ Thus $length(Q) = P(e)$. Why?

Proof of correctness of Dijkstra, continued

- ▶ Let Q be the (unique) path from s to y in T_{j+1} .
- ▶ **Observation.** The path Q contains a path from s to x in T_j and adds vertex y (and hence edge e) to form a path from s to y .
- ▶ Thus $\text{length}(Q) = P(e)$. Why?
- ▶ Now let R be *any* path in G from s to y .

Proof of correctness of Dijkstra, continued

- ▶ Let Q be the (unique) path from s to y in T_{j+1} .
- ▶ **Observation.** The path Q contains a path from s to x in T_j and adds vertex y (and hence edge e) to form a path from s to y .
- ▶ Thus $\text{length}(Q) = P(e)$. Why?
- ▶ Now let R be *any* path in G from s to y .
- ▶ To finish the proof, it suffices to show that $\text{length}(R) \geq \text{length}(Q)$. Why?

Proof of correctness of Dijkstra, continued




Figure: Green path is Q and orange + blue paths are R

Proof of correctness of Dijkstra, continued

- ▶ WLOG assume R and Q are different and let edge $f = vz$ be the first edge in path R that is not in T_j .

Proof of correctness of Dijkstra, continued

- ▶ WLOG assume R and Q are different and let edge $f = vz$ be the first edge in path R that is not in T_j .
- ▶ Assume also that $v \in V(T_j)$ (and hence $z \notin T_j$).

Proof of correctness of Dijkstra, continued

- ▶ WLOG assume R and Q are different and let edge $f = vz$ be the first edge in path R that is not in T_j .
- ▶ Assume also that $v \in V(T_j)$ (and hence $z \notin T_j$).
- ▶ Let K be the subpath of R from z to y .

Proof of correctness of Dijkstra, continued

- ▶ WLOG assume R and Q are different and let edge $f = vz$ be the first edge in path R that is not in T_j .
- ▶ Assume also that $v \in V(T_j)$ (and hence $z \notin T_j$).
- ▶ Let K be the subpath of R from z to y .
- ▶ Since e was a frontier edge used in the $(j + 1)$ st iteration of Algorithm Dijkstra, we have $P(e) \leq P(f)$. Why?

Dijkstra proof, continued: $P(e) \leq P(f)$

- ▶ In that case, $length(R) = dist[v] + \omega(f) + length(K)$

Dijkstra proof, continued: $P(e) \leq P(f)$

- ▶ In that case, $length(R) = dist[v] + \omega(f) + length(K)$
- ▶ $= P(f) + length(K)$

Dijkstra proof, continued: $P(e) \leq P(f)$

- ▶ In that case, $length(R) = dist[v] + \omega(f) + length(K)$
- ▶ $= P(f) + length(K)$
- ▶ $\geq P(e) = length(Q)$.

Dijkstra proof, continued: $P(e) \leq P(f)$

- ▶ In that case, $length(R) = dist[v] + \omega(f) + length(K)$
- ▶ $= P(f) + length(K)$
- ▶ $\geq P(e) = length(Q)$.
- ▶ In total, we have $length(R) \geq length(Q)$, as desired.

Dijkstra proof, continued: $P(e) \leq P(f)$

- ▶ In that case, $length(R) = dist[v] + \omega(f) + length(K)$
- ▶ $= P(f) + length(K)$
- ▶ $\geq P(e) = length(Q)$.
- ▶ In total, we have $length(R) \geq length(Q)$, as desired.
- ▶ **Conclusion.** For any $j \in \{0, 1, \dots, |V(G)| - 1\}$ we have that T_j is a shortest path tree for G . **QED**

A Question

Where did we use the fact that the edge weights of G are non-negative?

Final Remarks about Dijkstra

1. **Implementation.** Use a minimum priority queue by way of a binary heap: see section **24.3** for details.

Final Remarks about Dijkstra

1. **Implementation.** Use a minimum priority queue by way of a binary heap: see section **24.3** for details.
2. **BFS tree.** Algorithm Dijkstra gives us a method with which we can construct a **B**readth **F**irst **S**earch (BFS) tree.

BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.

BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.
- ▶ Choose a root vertex r and run G through Dijkstra's Algorithm.

BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.
- ▶ Choose a root vertex r and run G through Dijkstra's Algorithm.
- ▶ **Outcome.** In a Dijkstra tree,

BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.
- ▶ Choose a root vertex r and run G through Dijkstra's Algorithm.
- ▶ **Outcome.** In a Dijkstra tree,
 1. vertices that are distance **1** from r are at level **one**,


BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.
- ▶ Choose a root vertex r and run G through Dijkstra's Algorithm.
- ▶ **Outcome.** In a Dijkstra tree,
 1. vertices that are distance **1** from r are at level **one**,
 2. vertices that are distance **2** from r are at level **two**,

BFS Tree

- ▶ Assume the edges of graph G all have weight **1**.
- ▶ Choose a root vertex r and run G through Dijkstra's Algorithm.
- ▶ **Outcome.** In a Dijkstra tree,
 1. vertices that are distance **1** from r are at level **one**,
 2. vertices that are distance **2** from r are at level **two**,
 3. ETC

BFS tree, continued



Generalization of Dijkstra: Bellman-Ford Algorithm

- ▶ Recall that the input to Dijkstra's algorithms required that the edges are undirected and have non-negative weight.

Generalization of Dijkstra: Bellman-Ford Algorithm

- ▶ Recall that the input to Dijkstra's algorithms required that the edges are undirected and have non-negative weight.
- ▶ What happens if the weights are allowed to be from all of \mathbb{R} and are directed?

Generalization of Dijkstra: Bellman-Ford Algorithm

- ▶ Recall that the input to Dijkstra's algorithms required that the edges are undirected and have non-negative weight.
- ▶ What happens if the weights are allowed to be from all of \mathbb{R} and are directed?
- ▶ **Set-up.** Let G be a directed graph with edges weighted from \mathbb{R} and let s be the source vertex (i.e., the root).

Generalization of Dijkstra: Bellman-Ford Algorithm

- ▶ Recall that the input to Dijkstra's algorithms required that the edges are undirected and have non-negative weight.
- ▶ What happens if the weights are allowed to be from all of \mathbb{R} and are directed?
- ▶ **Set-up.** Let G be a directed graph with edges weighted from \mathbb{R} and let s be the source vertex (i.e., the root).
- ▶ Recall that $\omega(uv)$ is the weight of edge uv and since the edges are directed, order matters.

Generalization of Dijkstra: Bellman-Ford Algorithm

- ▶ Recall that the input to Dijkstra's algorithms required that the edges are undirected and have non-negative weight.
- ▶ What happens if the weights are allowed to be from all of \mathbb{R} and are directed?
- ▶ **Set-up.** Let G be a directed graph with edges weighted from \mathbb{R} and let s be the source vertex (i.e., the root).
- ▶ Recall that $\omega(uv)$ is the weight of edge uv and since the edges are directed, order matters.
- ▶ One of the main components of the **Bellman-Ford** algorithm is called **Edge Relaxation**.

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).
- ▶ Define the label $d[u]$ for each $u \in V(G)$, whose purpose is to give

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).
- ▶ Define the label $d[u]$ for each $u \in V(G)$, whose purpose is to give
- ▶ a **best approximation** for the distance from source vertex s to vertex u **so far**.

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).
- ▶ Define the label $d[u]$ for each $u \in V(G)$, whose purpose is to give
- ▶ a **best approximation** for the distance from source vertex s to vertex u **so far**.
- ▶ **Initialize** $d[u]$ to ∞ for every vertex $v \neq s$ and set $d[s] = 0$.

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).
- ▶ Define the label $d[u]$ for each $u \in V(G)$, whose purpose is to give
- ▶ a **best approximation** for the distance from source vertex s to vertex u **so far**.
- ▶ **Initialize** $d[u]$ to ∞ for every vertex $v \neq s$ and set $d[s] = 0$.
- ▶ We gather a collection of vertices C that ultimately become parts of paths from s to u ,

Edge Relaxation explained

- ▶ Let G be any weighted graph (directed or undirected).
- ▶ Define the label $d[u]$ for each $u \in V(G)$, whose purpose is to give
- ▶ a **best approximation** for the distance from source vertex s to vertex u **so far**.
- ▶ **Initialize** $d[u]$ to ∞ for every vertex $v \neq s$ and set $d[s] = 0$.
- ▶ We gather a collection of vertices C that ultimately become parts of paths from s to u ,
- ▶ beginning with the neighbors of s , and then the neighbors of the neighbors of s , and so on.

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,
- ▶ and then add u to the set C ;

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,
- ▶ and then add u to the set C ;
- ▶ then update the label $d[z]$ for all vertices z that are neighbors of u and that are not already in C .

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,
- ▶ and then add u to the set C ;
- ▶ then update the label $d[z]$ for all vertices z that are neighbors of u and that are not already in C .
- ▶ The procedure so described allows the algorithm to detect if there is a shorter way to get from s to u by way of z .

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,
- ▶ and then add u to the set C ;
- ▶ then update the label $d[z]$ for all vertices z that are neighbors of u and that are not already in C .
- ▶ The procedure so described allows the algorithm to detect if there is a shorter way to get from s to u by way of z .
- ▶ This technique is called **edge relaxation**: it takes an old estimate of the distance from s to u , and

Edge relaxation, continued

- ▶ In particular, at each iteration of the algorithm, we choose a vertex $u \notin C$ with the smallest possible $d[u]$ label,
- ▶ and then add u to the set C ;
- ▶ then update the label $d[z]$ for all vertices z that are neighbors of u and that are not already in C .
- ▶ The procedure so described allows the algorithm to detect if there is a shorter way to get from s to u by way of z .
- ▶ This technique is called **edge relaxation**: it takes an old estimate of the distance from s to u , and
- ▶ and tries to improve (ie, decrease) the value of the distance label $d[u]$.

Pseudocode for Edge Relaxation

► **if** $d[u] + \omega(uz) < d[z]$ **then**

Pseudocode for Edge Relaxation


- ▶ **if** $d[u] + \omega(uz) < d[z]$ **then**
- ▶ $d[z] = d[u] + \omega(uz)$

Reachable negative-weight cycles are bad

- ▶ Before giving the Bellman-Ford algorithm, let's look at scenario in a graph for which there is no shortest path tree.

Reachable negative-weight cycles are bad

- Before giving the Bellman-Ford algorithm, let's look at scenario in a graph for which there is no shortest path tree.



Negative-weight cycles are bad

- ▶ **The Point.** If there exists a negative-weight cycle reachable from source vertex s then there is no shortest path from s to any other vertex reachable from s .

Bellman-Ford Algorithm (**B-F**(G, s))

- **Input.** G , a directed graph and $s \in V(G)$, a source vertex.

Bellman-Ford Algorithm (**B-F**(G, s))

- ▶ **Input.** G , a directed graph and $s \in V(G)$, a source vertex.
- ▶ **Output.** Either **TRUE** if there is no reachable-from- s negative weight cycle or **FALSE**. The former will also keep track of the shortest path tree.

B-F(G, s), continued

- ▶ $d[v] = \infty$ for all $v \in V(G)$
- ▶ $d[s] = 0$
- ▶ **For** $i = 1$ **to** $|V(G)| - 1$
- ▶ **For** each $uv \in E(G)$
- ▶ **if** $d[u] + \omega(uv) < d[v]$
- ▶ **then** $d[v] = d[u] + \omega(uv)$
- ▶ $parent[v] = u$
- ▶ **For** each edge uv
- ▶ **if** $d[u] + \omega(uv) < d[v]$
- ▶ **then return** “False”
- ▶ **return** “True”

The **red code** is the edge-relaxation part of the algorithm.

Some Analysis of Bellman-Ford

- ▶ For a full explanation of the correctness of Bellman-Ford, see chapter 24 in our text.

Some Analysis of Bellman-Ford

- ▶ For a full explanation of the correctness of Bellman-Ford, see chapter 24 in our text.
- ▶ We'll zoom in on the line

Some Analysis of Bellman-Ford

- ▶ For a full explanation of the correctness of Bellman-Ford, see chapter 24 in our text.
- ▶ We'll zoom in on the line
- ▶ **if $d[u] + \omega(uv) < d[v]$ then return "False"**

Some Analysis of Bellman-Ford

- ▶ For a full explanation of the correctness of Bellman-Ford, see chapter 24 in our text.
- ▶ We'll zoom in on the line
- ▶ **if $d[u] + \omega(uv) < d[v]$ then return "False"**
- ▶ and prove the following claim:

Some Analysis of Bellman-Ford

- ▶ For a full explanation of the correctness of Bellman-Ford, see chapter 24 in our text.
- ▶ We'll zoom in on the line
- ▶ **if $d[u] + \omega(uv) < d[v]$ then return "False"**
- ▶ and prove the following claim:
- ▶ **Claim.** If the input graph G contains a negative weight cycle reachable from s , then **B-F**(G, s) returns "False."

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”
- ▶ Then $d[v_i] + \omega(v_i v_{i+1}) \geq d[v_{i+1}] \quad \forall i \in \{0, 1, \dots, k-1\}$.

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”
- ▶ Then $d[v_i] + \omega(v_i v_{i+1}) \geq d[v_{i+1}] \quad \forall i \in \{0, 1, \dots, k-1\}$.
- ▶ $\Rightarrow \sum_{i=0}^{k-1} d[v_i] + \sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq \sum_{i=0}^{k-1} d[v_{i+1}]$.

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”
- ▶ Then $d[v_i] + \omega(v_i v_{i+1}) \geq d[v_{i+1}] \quad \forall i \in \{0, 1, \dots, k-1\}$.
- ▶ $\Rightarrow \sum_{i=0}^{k-1} d[v_i] + \sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq \sum_{i=0}^{k-1} d[v_{i+1}]$.
- ▶ Observe that $\sum_{i=0}^{k-1} d[v_i] = \sum_{i=0}^{k-1} d[v_{i+1}]$. Why?

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”
- ▶ Then $d[v_i] + \omega(v_i v_{i+1}) \geq d[v_{i+1}] \quad \forall i \in \{0, 1, \dots, k-1\}$.
- ▶ $\Rightarrow \sum_{i=0}^{k-1} d[v_i] + \sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq \sum_{i=0}^{k-1} d[v_{i+1}]$.
- ▶ Observe that $\sum_{i=0}^{k-1} d[v_i] = \sum_{i=0}^{k-1} d[v_{i+1}]$. Why?
- ▶ Thus $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq 0$, which is the desired contradiction.

Proof of Claim

- ▶ Let $v_0 v_1 \dots v_k$ with $v_0 = v_k$ be a negative weight cycle.
- ▶ That is $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) < 0$.
- ▶ BWOC assume that algorithm **B-F** does not output “False.”
- ▶ Then $d[v_i] + \omega(v_i v_{i+1}) \geq d[v_{i+1}] \quad \forall i \in \{0, 1, \dots, k-1\}$.
- ▶ $\Rightarrow \sum_{i=0}^{k-1} d[v_i] + \sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq \sum_{i=0}^{k-1} d[v_{i+1}]$.
- ▶ Observe that $\sum_{i=0}^{k-1} d[v_i] = \sum_{i=0}^{k-1} d[v_{i+1}]$. Why?
- ▶ Thus $\sum_{i=0}^{k-1} \omega(v_i v_{i+1}) \geq 0$, which is the desired contradiction.
- ▶ **QED**

Runtime of **B-F**

- ▶ Assume the input graph G is simple.

Runtime of **B-F**

- ▶ Assume the input graph G is simple.
- ▶ Then the runtime of algorithm **B-F** is $O(|V(G)|^3)$.

Runtime of **B-F**

- ▶ Assume the input graph G is simple.
- ▶ Then the runtime of algorithm **B-F** is $O(|V(G)|^3)$.
- ▶ **Proof:** Exercise.

Runtime of **B-F**

- ▶ Assume the input graph G is simple.
- ▶ Then the runtime of algorithm **B-F** is $O(|V(G)|^3)$.
- ▶ **Proof:** Exercise.
- ▶ Finally, a good alternative source for both Dijkstra and Bellman-Ford is Gross and Yellen's textbook **Graph Theory and Applications**, especially for the proof of Dijkstra and the use of frontier edges in tree growing.

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.
- ▶ A **flow network** is a directed graph G in which every edge $uv \in E(G)$ has a positive **capacity** $c(uv)$, and

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.
- ▶ A **flow network** is a directed graph G in which every edge $uv \in E(G)$ has a positive **capacity** $c(uv)$, and
- ▶ furthermore $c(uv) = 0$ if uv is not an edge.

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.
- ▶ A **flow network** is a directed graph G in which every edge $uv \in E(G)$ has a positive **capacity** $c(uv)$, and
- ▶ furthermore $c(uv) = 0$ if uv is not an edge.
- ▶ In other words, the function c describes the weights on the edges and non-edges of G .

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.
- ▶ A **flow network** is a directed graph G in which every edge $uv \in E(G)$ has a positive **capacity** $c(uv)$, and
- ▶ furthermore $c(uv) = 0$ if uv is not an edge.
- ▶ In other words, the function c describes the weights on the edges and non-edges of G .
- ▶ There is a **source** vertex s and a **sink** vertex t .

New Topic in Graph Algorithms: Network Flows

- ▶ We'll begin with some definitions.
- ▶ A **flow network** is a directed graph G in which every edge $uv \in E(G)$ has a positive **capacity** $c(uv)$, and
- ▶ furthermore $c(uv) = 0$ if uv is not an edge.
- ▶ In other words, the function c describes the weights on the edges and non-edges of G .
- ▶ There is a **source** vertex s and a **sink** vertex t .
- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following properties:

Network flow definitions and constraints, continued

- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following **three** properties:

Network flow definitions and constraints, continued

- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following **three** properties:
 1. **Capacity Constraint:** $f(uv) \leq c(uv) \quad \forall u, v \in V(G)$. “can't burst the pipes”

Network flow definitions and constraints, continued

- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following **three** properties:
 1. **Capacity Constraint:** $f(uv) \leq c(uv) \quad \forall u, v \in V(G)$. “can't burst the pipes”
 2. **Skew Symmetry:** $f(uv) = -f(vu)$. “positive flow in one direction = negative flow in the opposite direction.”

Network flow definitions and constraints, continued

- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following **three** properties:
 1. **Capacity Constraint:** $f(uv) \leq c(uv) \quad \forall u, v \in V(G)$. “can't burst the pipes”
 2. **Skew Symmetry:** $f(uv) = -f(vu)$. “positive flow in one direction = negative flow in the opposite direction.”
 3. **Flow Conservation:** $\sum_{v \in V(G)} f(uv) = 0 \quad \forall v \in V(G) \setminus \{s, t\}$. “amount of water entering u = amount of water leaving u ”

Network flow definitions and constraints, continued


- ▶ A **flow** is a real-valued function $f : V \times V \rightarrow \mathbb{R}$ that has the following **three** properties:
 1. **Capacity Constraint:** $f(uv) \leq c(uv) \quad \forall u, v \in V(G)$. “can’t burst the pipes”
 2. **Skew Symmetry:** $f(uv) = -f(vu)$. “positive flow in one direction = negative flow in the opposite direction.”
 3. **Flow Conservation:** $\sum_{v \in V(G)} f(uv) = 0 \quad \forall v \in V(G) \setminus \{s, t\}$. “amount of water entering u = amount of water leaving u ”
- ▶ Finally, the **value of flow** f is $\sum_{v \in V(G) \setminus s} f(sv)$. That is, the flow is the amount of water leaving the source vertex s .

Example from Herb Wilf's downloadable Algorithms text

- ▶ The network flow example below is from Chapter 3 of <https://www.math.upenn.edu/~wilf/AlgComp3.html>

Example from Herb Wilf's downloadable Algorithms text


- ▶ The network flow example below is from Chapter 3 of <https://www.math.upenn.edu/~wilf/AlgComp3.html>



- ▶ The number in the \square is the flow and can change.
- ▶ The number in the non-box is the capacity and can't change.

Example from Herb Wilf's downloadable Algorithms text


- ▶ The network flow example below is from Chapter 3 of <https://www.math.upenn.edu/~wilf/AlgComp3.html>



- ▶ The number in the \square is the flow and can change.
- ▶ The number in the non-box is the capacity and can't change.
- ▶ **Exercise.** Check that this network satisfies skew symmetry, the capacity constraint, and flow conservation.


Example from Herb Wilf's downloadable Algorithms text

- ▶ The network flow example below is from Chapter 3 of <https://www.math.upenn.edu/~wilf/AlgComp3.html>




- ▶ The number in the \square is the flow and can change.
- ▶ The number in the non-box is the capacity and can't change.
- ▶ **Exercise.** Check that this network satisfies skew symmetry, the capacity constraint, and flow conservation.

Example from Herb Wilf's downloadable Algorithms text



Example from Herb Wilf's downloadable Algorithms text



- **Exercise.** What is the flow of the above network?

Hmmmm

- ▶ We haven't stated a problem yet!!

Hmmmm

- ▶ We haven't stated a problem yet!!
- ▶ **Here goes.**

Hmmmm

- ▶ We haven't stated a problem yet!!
- ▶ **Here goes.**
- ▶ **The Network Flow Problem.** Given a flow network G with capacity constraint function c , find the flow of maximum value.

Source for applications

- ▶ For one, of many, interesting sources for applications of the network flow problem, see

Source for applications

- ▶ For one, of many, interesting sources for applications of the network flow problem, see
- ▶ `http://blogs.cornell.edu/info2040/2012/11/04/applications-of-network-flow/`

Source for applications

- ▶ For one, of many, interesting sources for applications of the network flow problem, see
- ▶ <http://blogs.cornell.edu/info2040/2012/11/04/applications-of-network-flow/>
- ▶ **Next Week.** More on network flows by way of a side-trip to linear programming, the max-flow/min cut theorem, and the Ford-Fulkerson hall of fame.